

## 3.2.1)

Imbench\_results (~/.pa2-gornad/Data) - gedit

File Edit View Search Tools Documents Help

Open Save Undo

Imbench\_results

1 Context switching - times in microseconds - smaller is better

2 -----

Host	OS	2p/0K ctxsw	2p/16K ctxsw	2p/64K ctxsw	8p/16K ctxsw	8p/64K ctxsw	16p/16K ctxsw	16p/64K ctxsw
cse-p204i	Linux 2.6.32-	6.4100	7.4200	11.8	9.7300	15.1	10.5	15.2
cse-p204i	Linux 2.6.32-	6.4600	5.9100	12.4	9.1600	13.8	9.6200	15.0
cse-p204i	Linux 2.6.32-	7.3000	9.1600	7.4000	10.4	14.3	10.3	14.7
cse-p204i	Linux 2.6.32-	7.3000	7.8400	13.5	10.4	15.4	10.8	15.8
cse-p204i	Linux 2.6.32-	7.8800	8.1300	13.1	10.3	15.7	10.5	15.2

11

Empirical average and variance of kernel threads performance:

Average =  $(6.4100 + 6.4600 + 7.3000 + 7.3000 + 7.8800) / 5 = 7.07 \mu s$

Variance = 0.3924 (using excel)

The number of context switches so far is: 1

The time taken by the context switch is: 2025 = 2.025  $\mu s$

The number of context switches so far is: 2

The time taken by the context switch is: 2095 = 2.095  $\mu s$

The number of context switches so far is: 3

The time taken by the context switch is: 1886 = 1.886  $\mu s$

The number of context switches so far is: 4

The time taken by the context switch is: 2374 = 2.374  $\mu s$

The number of context switches so far is: 5

The time taken by the context switch is: 2095 = 2.095  $\mu s$

The number of context switches so far is: 6

The time taken by the context switch is: 2165 = 2.165  $\mu s$

The number of context switches so far is: 7

The time taken by the context switch is: 1886 = 1.886  $\mu s$

The number of context switches so far is: 8

The time taken by the context switch is: 2025 = 2.025  $\mu s$

The number of context switches so far is: 9

The time taken by the context switch is: 1816 = 1.816  $\mu s$

The number of context switches so far is: 10

The time taken by the context switch is: 2025 = 2.025  $\mu s$

The number of context switches so far is: 11

The time taken by the context switch is:  $1257 = 1.257 \mu\text{s}$

Empirical average and variance:

Average =  $(2025 + 2095 + 1886 + 2374 + 2095 + 2165 + 1886 + 2025 + 1816 + 2025 + 1257)/11 = 1968 \text{ ns} = 1.968 \mu\text{s}$

Variance = 78772.29 (using excel)

3.2.2)

Kernel Threads Sleeping()	User Threads Sleeping()
17463847	17632545
16270248	19306493
17177640	22372854
17669386	28898776
17209392	43898551
17069643	72135502




Image 1: Kernel threads sleeping()




Image 2: User threads sleeping()

Note: Both user and kernel thread tests were tested for an increasing number of threads similar to 3.2.3 to produce more reliable results.

For kernel threads, we notice that performance is extremely unreliable. This is caused by the expensive context switching times coupled with occupied system resources and scheduling. For user level threads, we notice a consistent increase of performance that is due to the relatively inexpensive context switching times of user threads.

### 3.2.3)

Kernel Threads Counter()	User Threads Counter()
17329426	17656769
18020844	18867391
17065985	22594306
16195997	29240851
16545564	43117937
17784652	71651928




Image 3: Kernel threads counter()

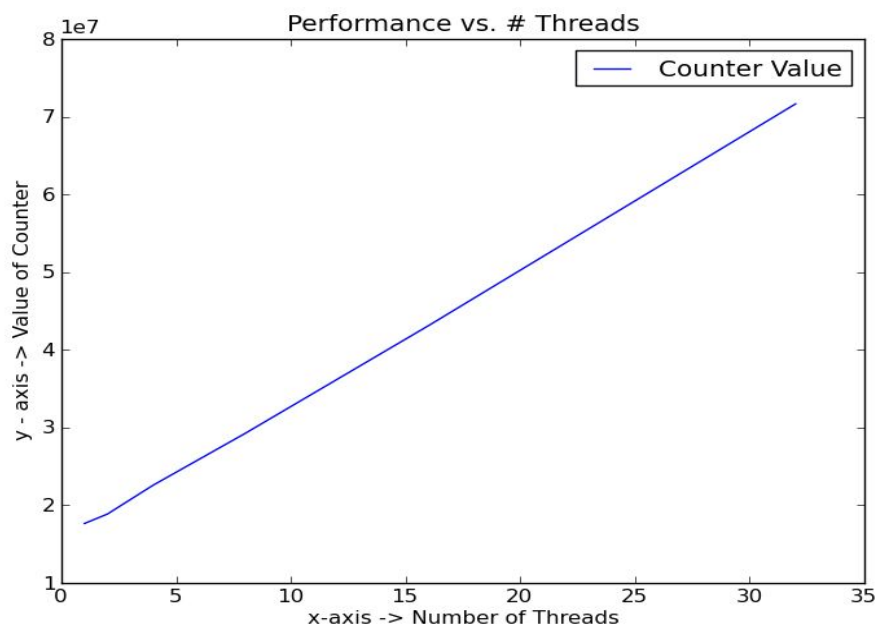


Image 4: User threads counter()

Since the counter() function is less expensive for kernel threads, we do not see as much inconsistency as with the sleeping() function. For both graphs, we notice an increase of work done in relation to the number of threads. There is a correlation between the number of cores in the system and the number of threads. This is explained by parallel processing. More cores mean more threads can be handled efficiently leading to an increased performance.